



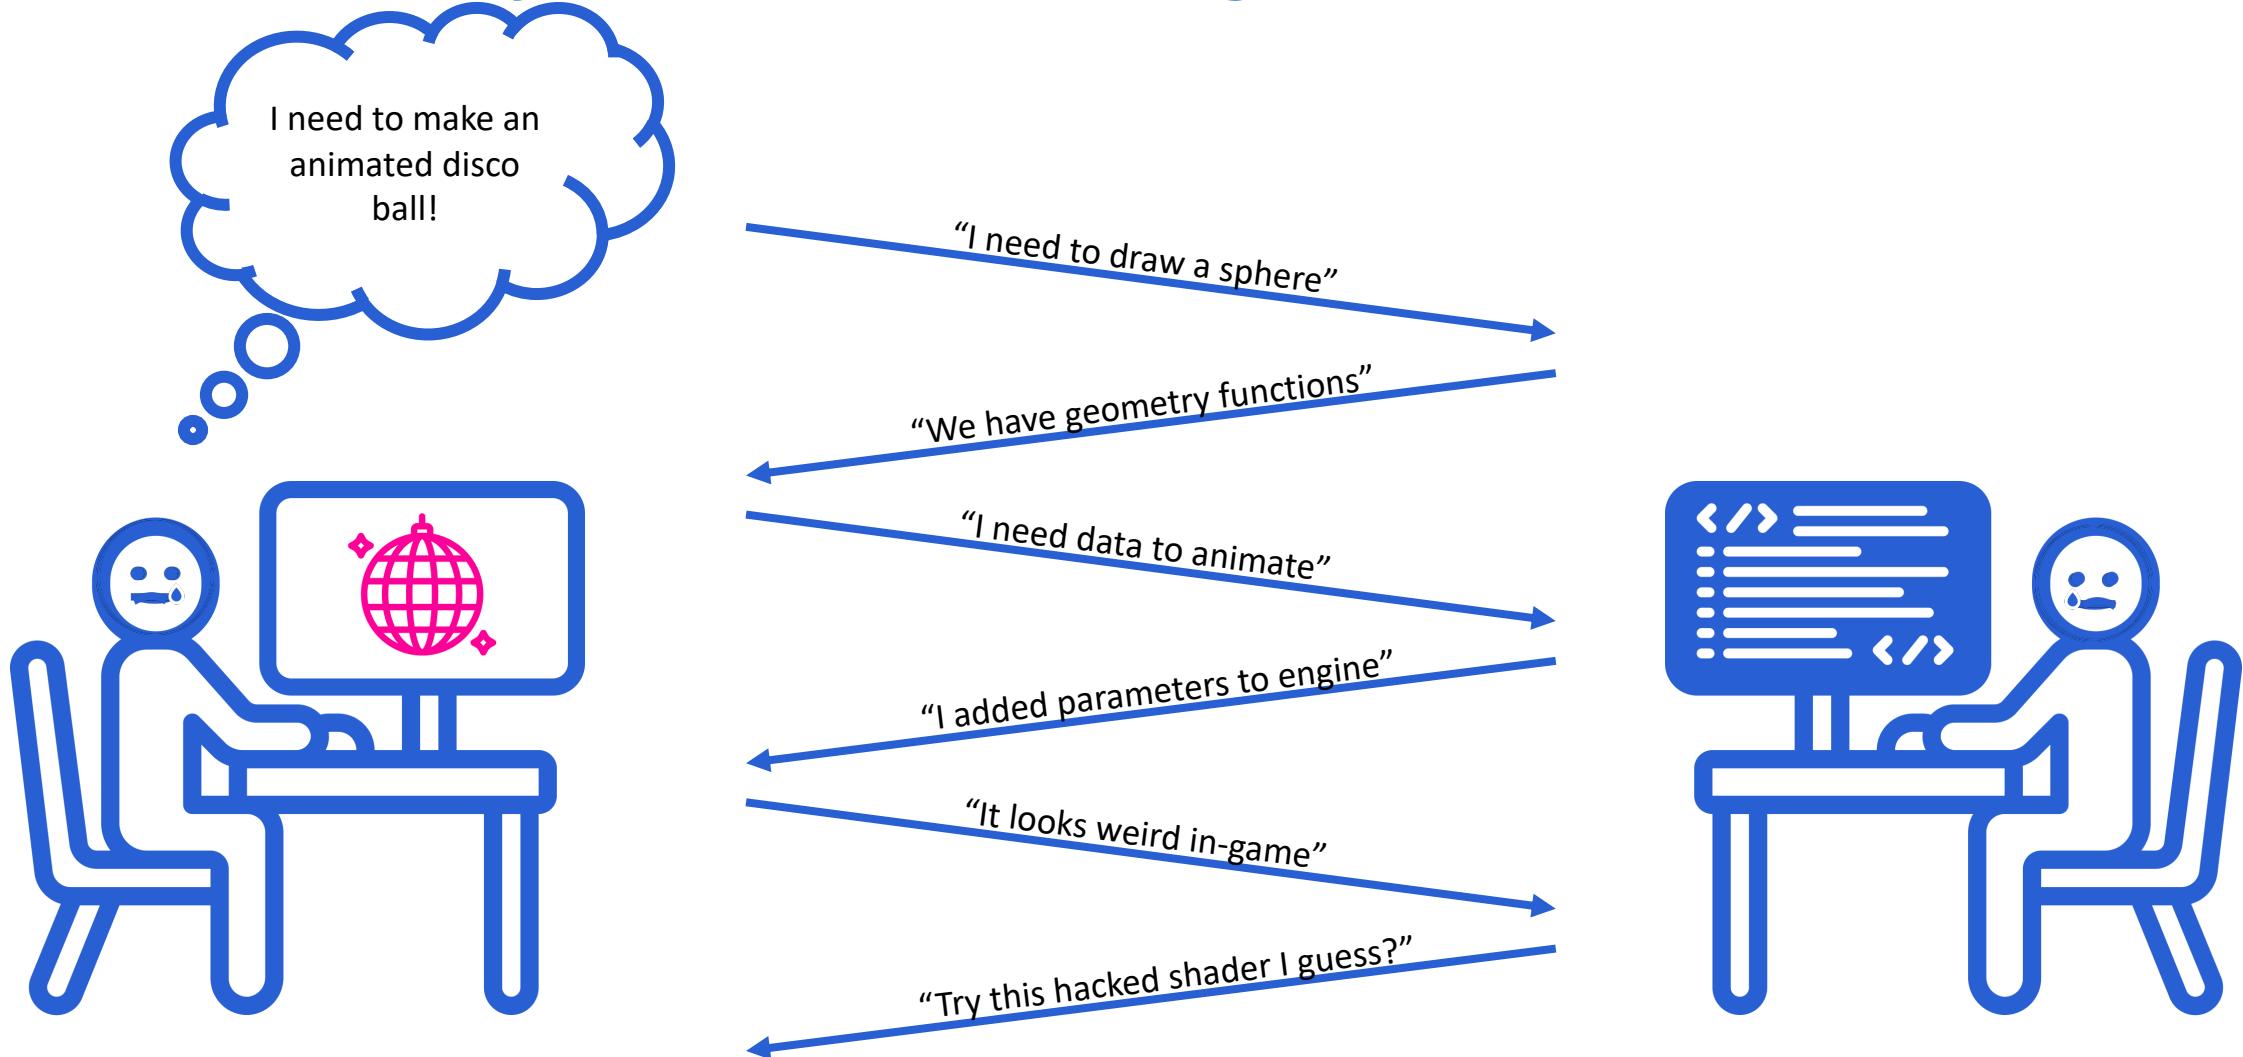
GPUs Unleashed!

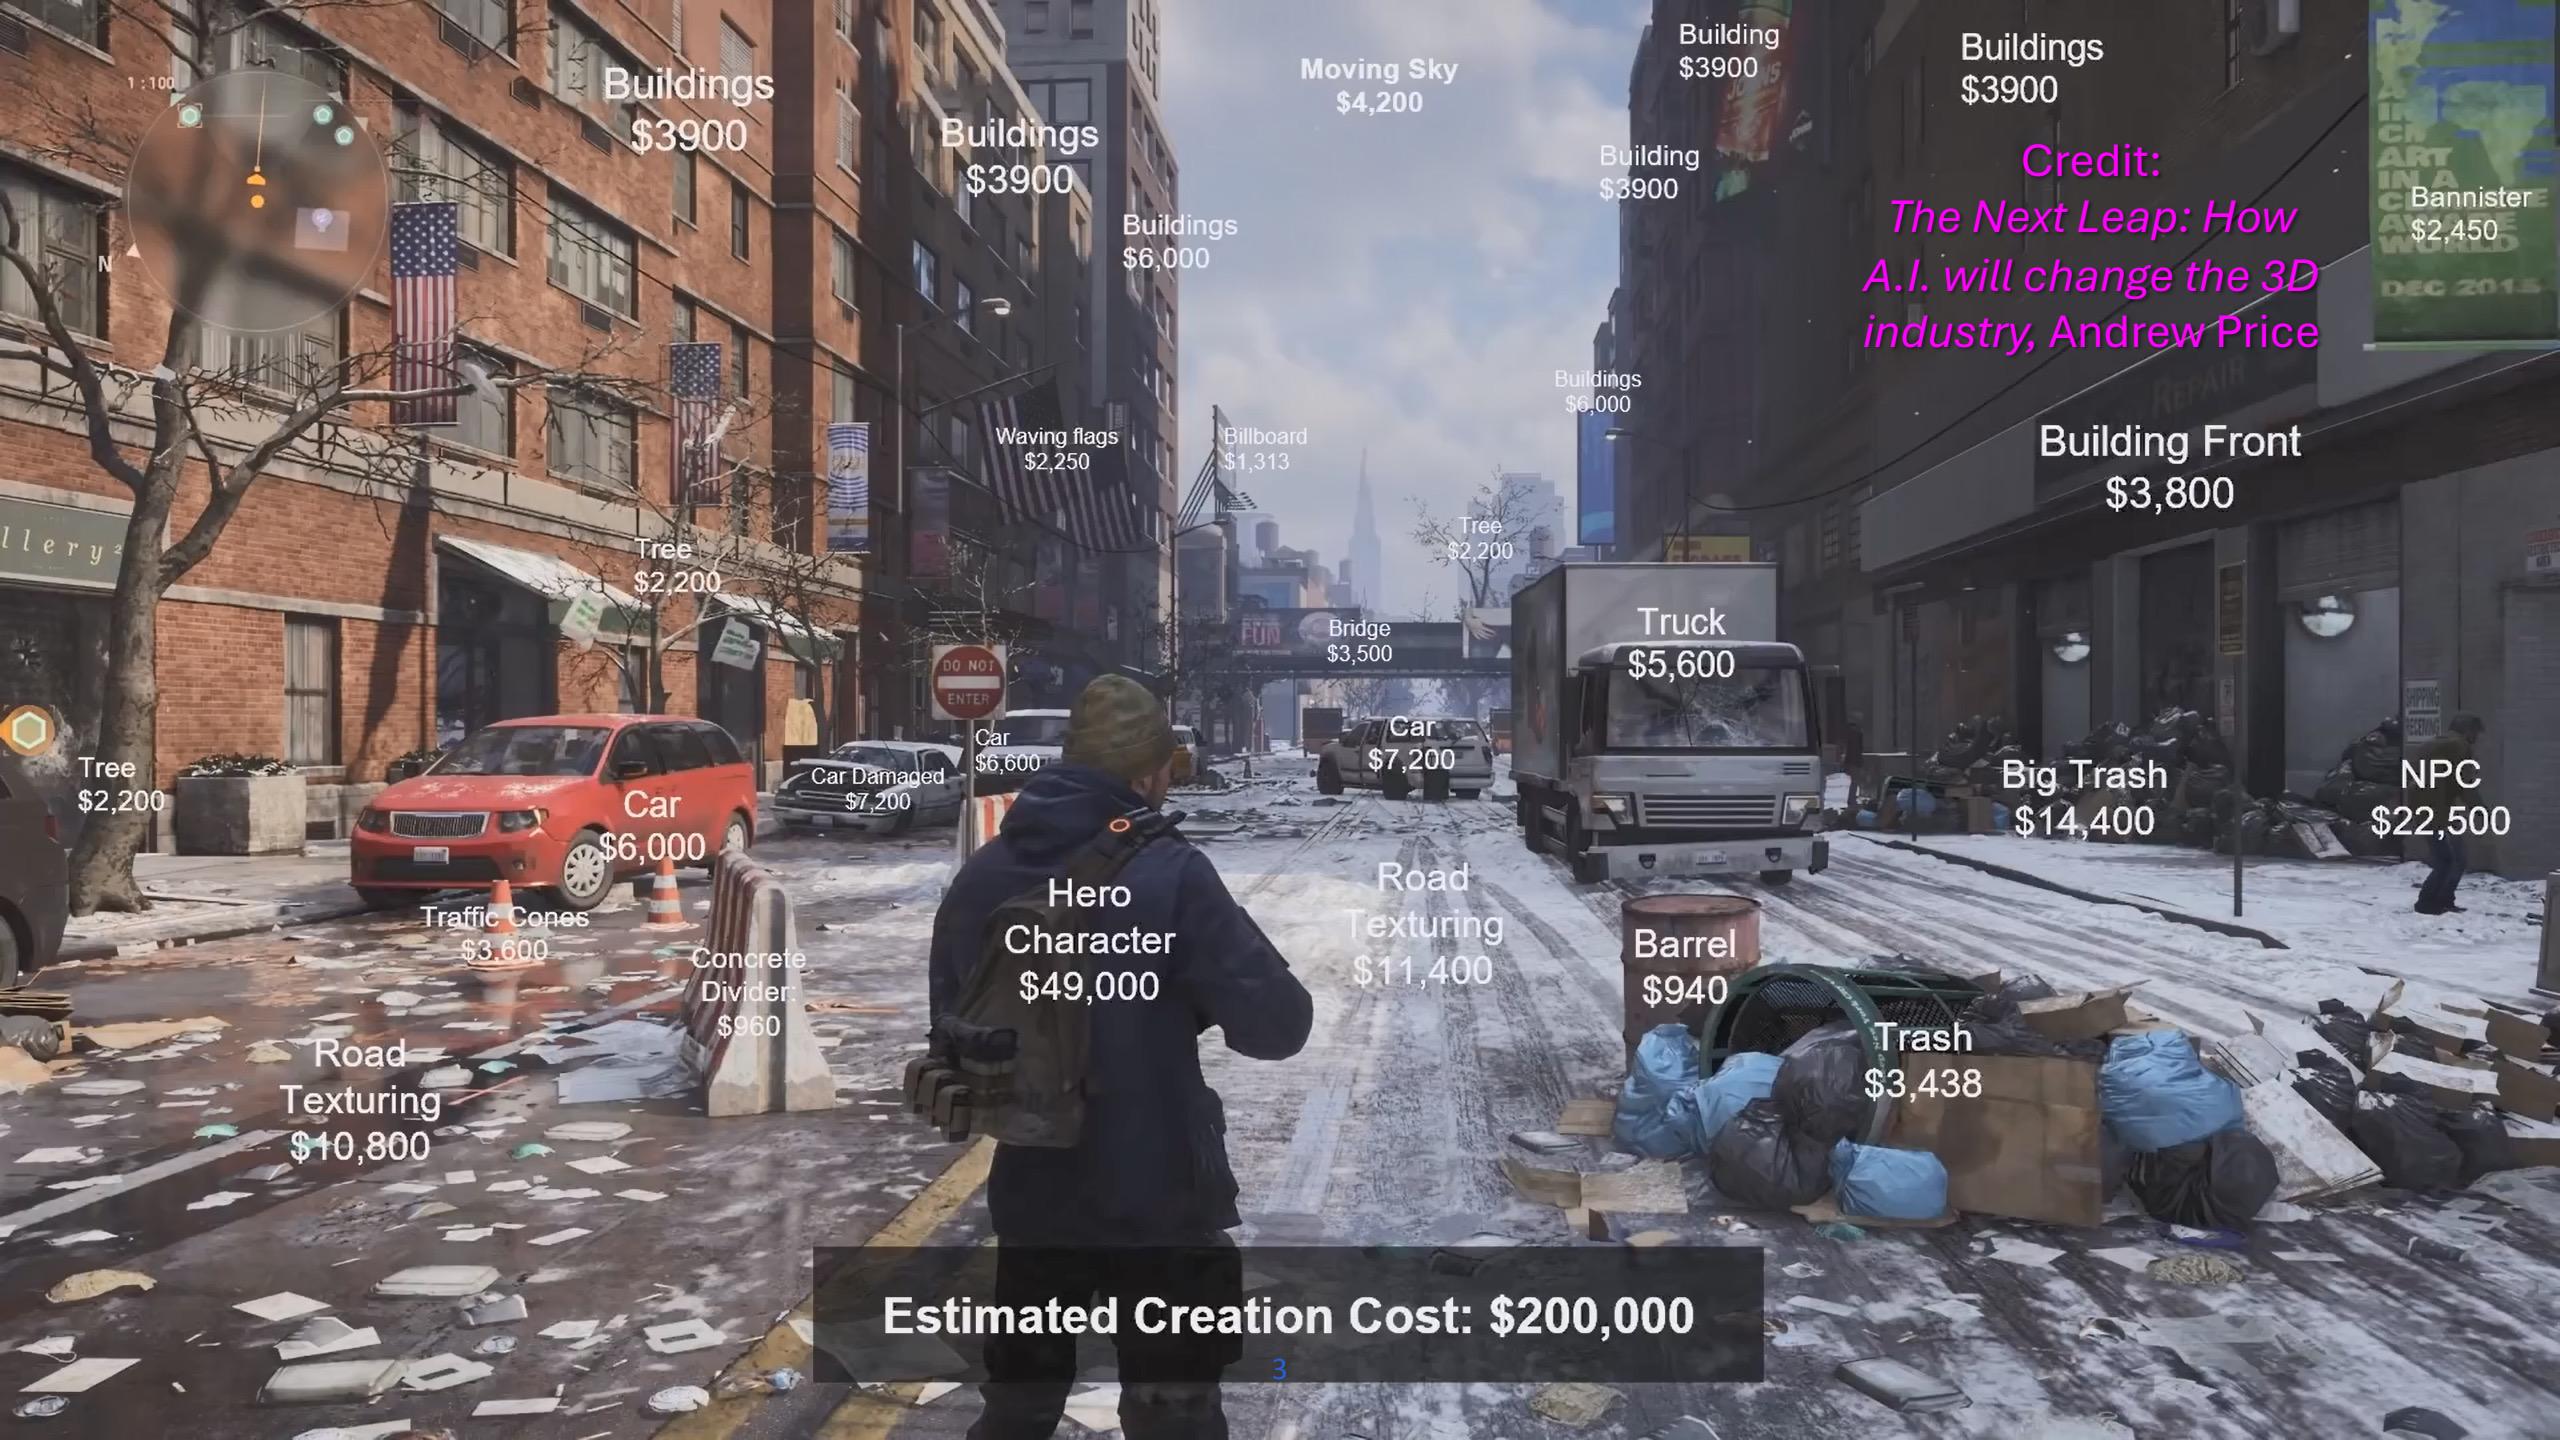
Make Your Games More Powerful With wasi-gfx

Sean Isom, CEO @ Renderlet

Mendy Berger, Sr. Software Engineer @ Renderlet

The bits and bytes of 3D scenes in games





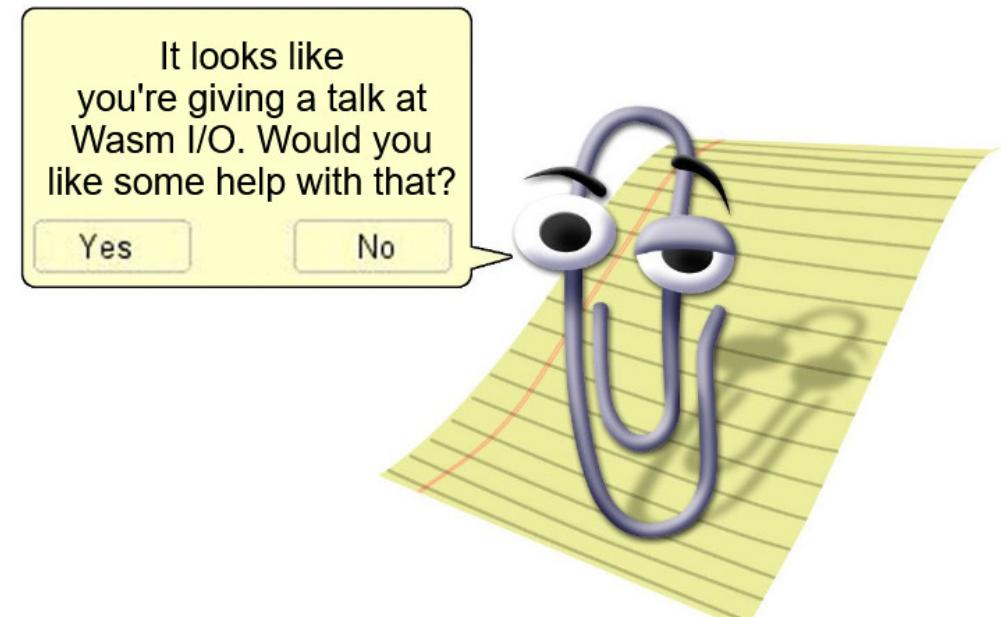
The big problem: content

- Content is expensive to build!
- Exploding scope, quantity, and quality bar
- Data is intertwined with code for 3D
- Not necessarily reusable
- GLTF is for static data



Plugins to the rescue

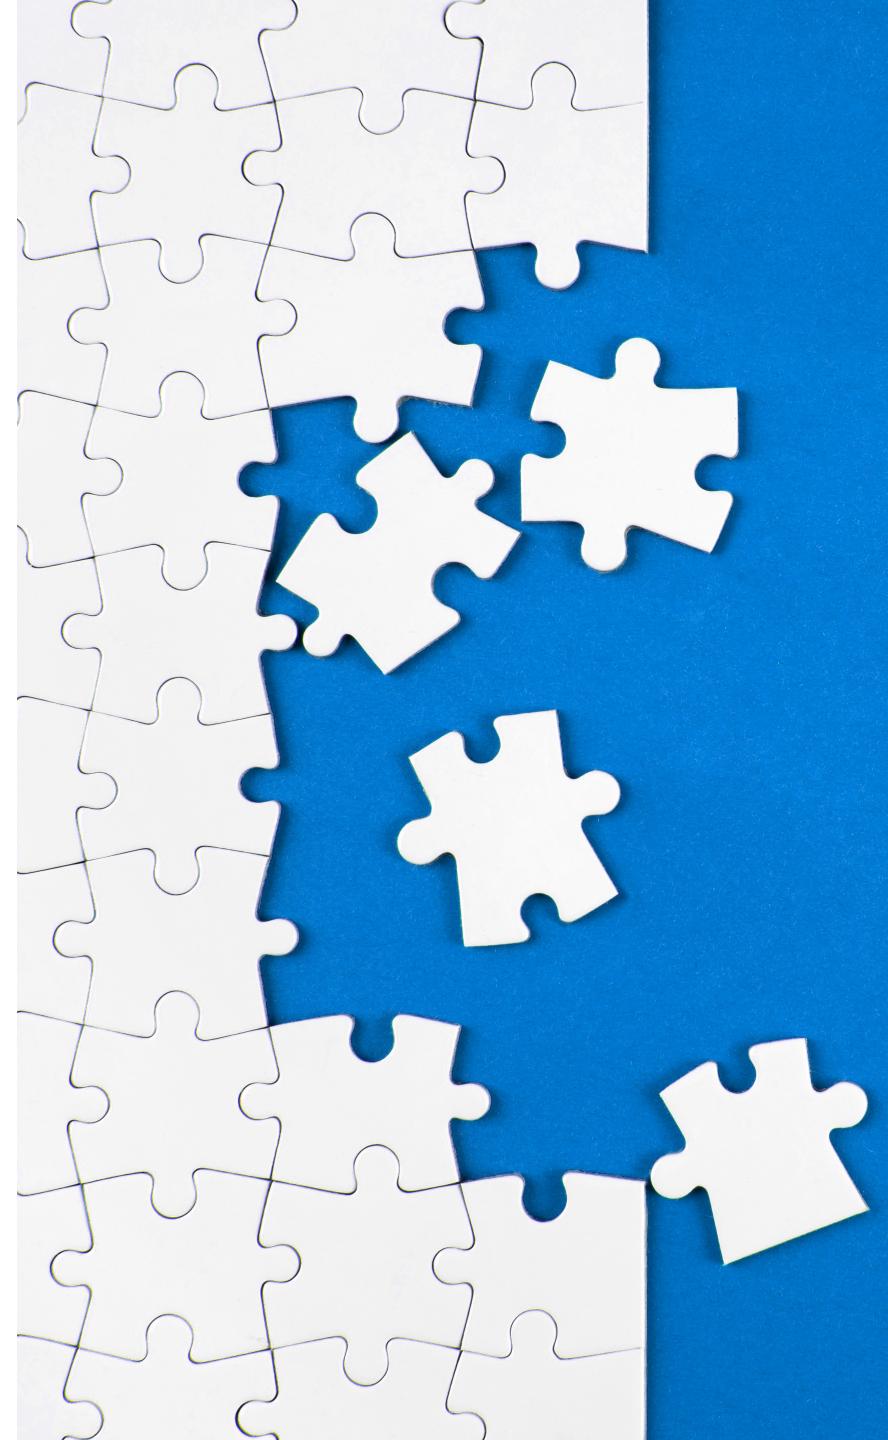
- Rendering becomes plugin-driven
- Definition: 3rd party code called over a defined interface
- Plugins come in many shapes and sizes
 - WinAmp
 - Photoshop
 - Kubernetes CRDs
 - Browser APIs
 - Clippy
- Games are different! Plugins need a GPU



imgflip.com

Wasm + GPUs = ?

- WebGPU: multitenant GPU security
- Wasm: a true native sandbox (not Flash)
- In-browser:
 - Browser GPU APIs over JS Bridge
 - Emscripten / wasm-bindgen
- Outside browser:
 - ?
 - Wasm Component Model to the rescue!

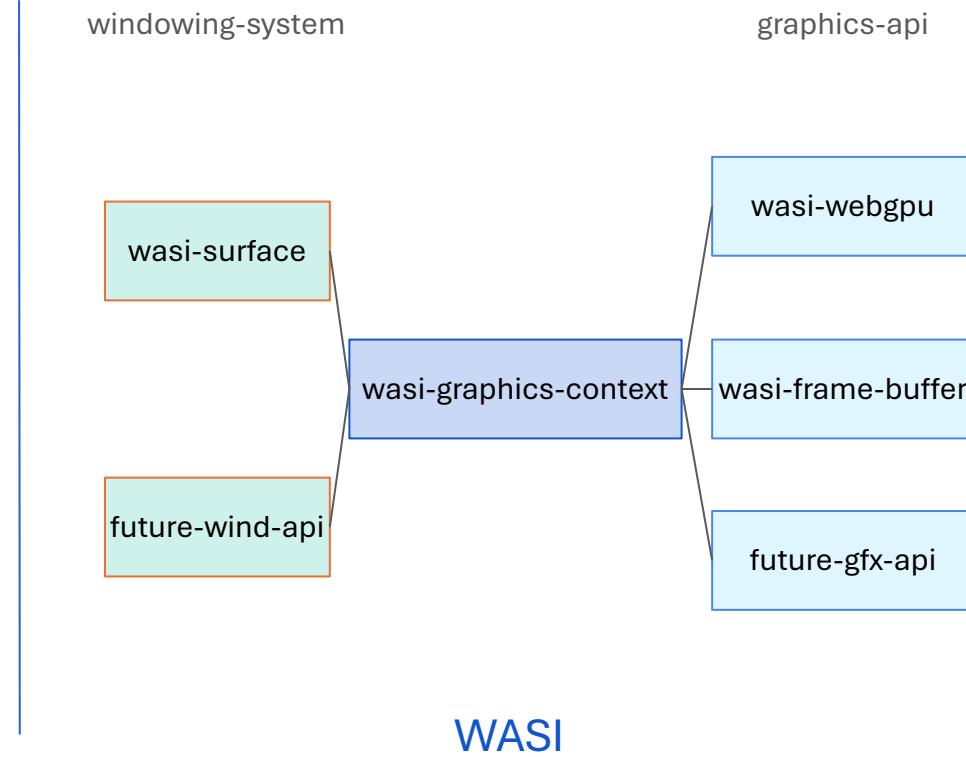
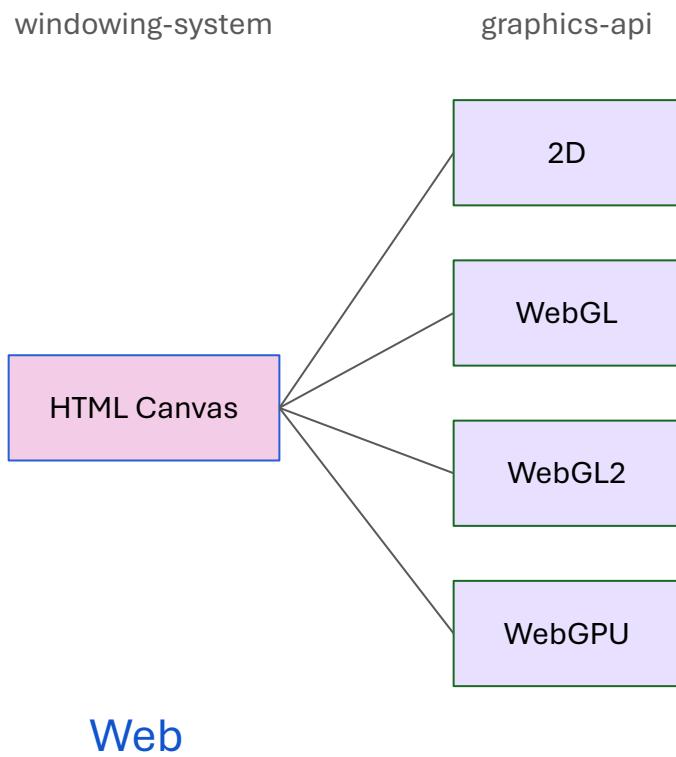


Introducing wasi-gfx

- Machine translate WebGPU's WebIDL to wit
- Extend with support to connect with surfaces and IO
- Window and graphics API agnostic
- Reference host implementation: wasi-gfx-runtime
 - <https://github.com/wasi-gfx/wasi-gfx-runtime>

Demo: Bevy

wasi-gfx architecture



Considerations for drawing into a window

- “Hooking” backbuffer is straightforward
 - Consider subregions for “applets” like UI
 - Fullscreen quad / not copy
- Drawing into scene in 3D is difficult
 - Host context – deferred targets, screen-space effects, etc.
 - Shared camera/materials between host and guest
 - Consider rendering into own texture in 3D with alpha
 - Also consider using host app’s shader on plugin data

Demo: Maps

Using wasi-gfx in a game

- Like a box of “Lego” parts
- Need a plugin runtime
 - Renderlet is building this
- Adoption hurdles:
 - Games are C++
 - Games demand performance
 - Games aren’t WGPU



C++/DX12

Rust/WGPU



Hurdle: Integrating C++ and Rust

- Plugins written in any language
- Host for wasi-gfx requires Rust*
 - WGPU
 - Wasmtime (Component model)
 - wasi-gfx-runtime
- Maintain separate build tree
 - C ABI “DLL”
 - Bindings via cbindgen / cxx
 - Lots of #[repr(C)] and void*
 - Clean C++ API wrapper

Hurdle: Real-time performance

- Wasm usually is fast on CPU (~80%)
- Wasm shouldn't be the bottleneck
- Manage resources intentionally
 - Don't readback data to CPU
 - Move to GPU-driven rendering
 - Avoid allocs with memory mapping

Hurdle: Integrating DX12 and WebGPU

- WebGPU implemented through WGPU in wasi-gfx-runtime
- WGPU uses DX12 backend on Windows
- Change WGPU to accept device created by host
- “as_hal” -> Get guest-created WGPU resource as DX12 for host
- “from_hal” -> Wrap host-created DX12 resource for WGPU guest
- Only required to share small number of resources



Introducing:
Renderlet Plugin Runtime



Demo: Adria

Next: Sane data models for plugins



Don't

- Export entire host app API (headers)
- Build bespoke, one-off mappings

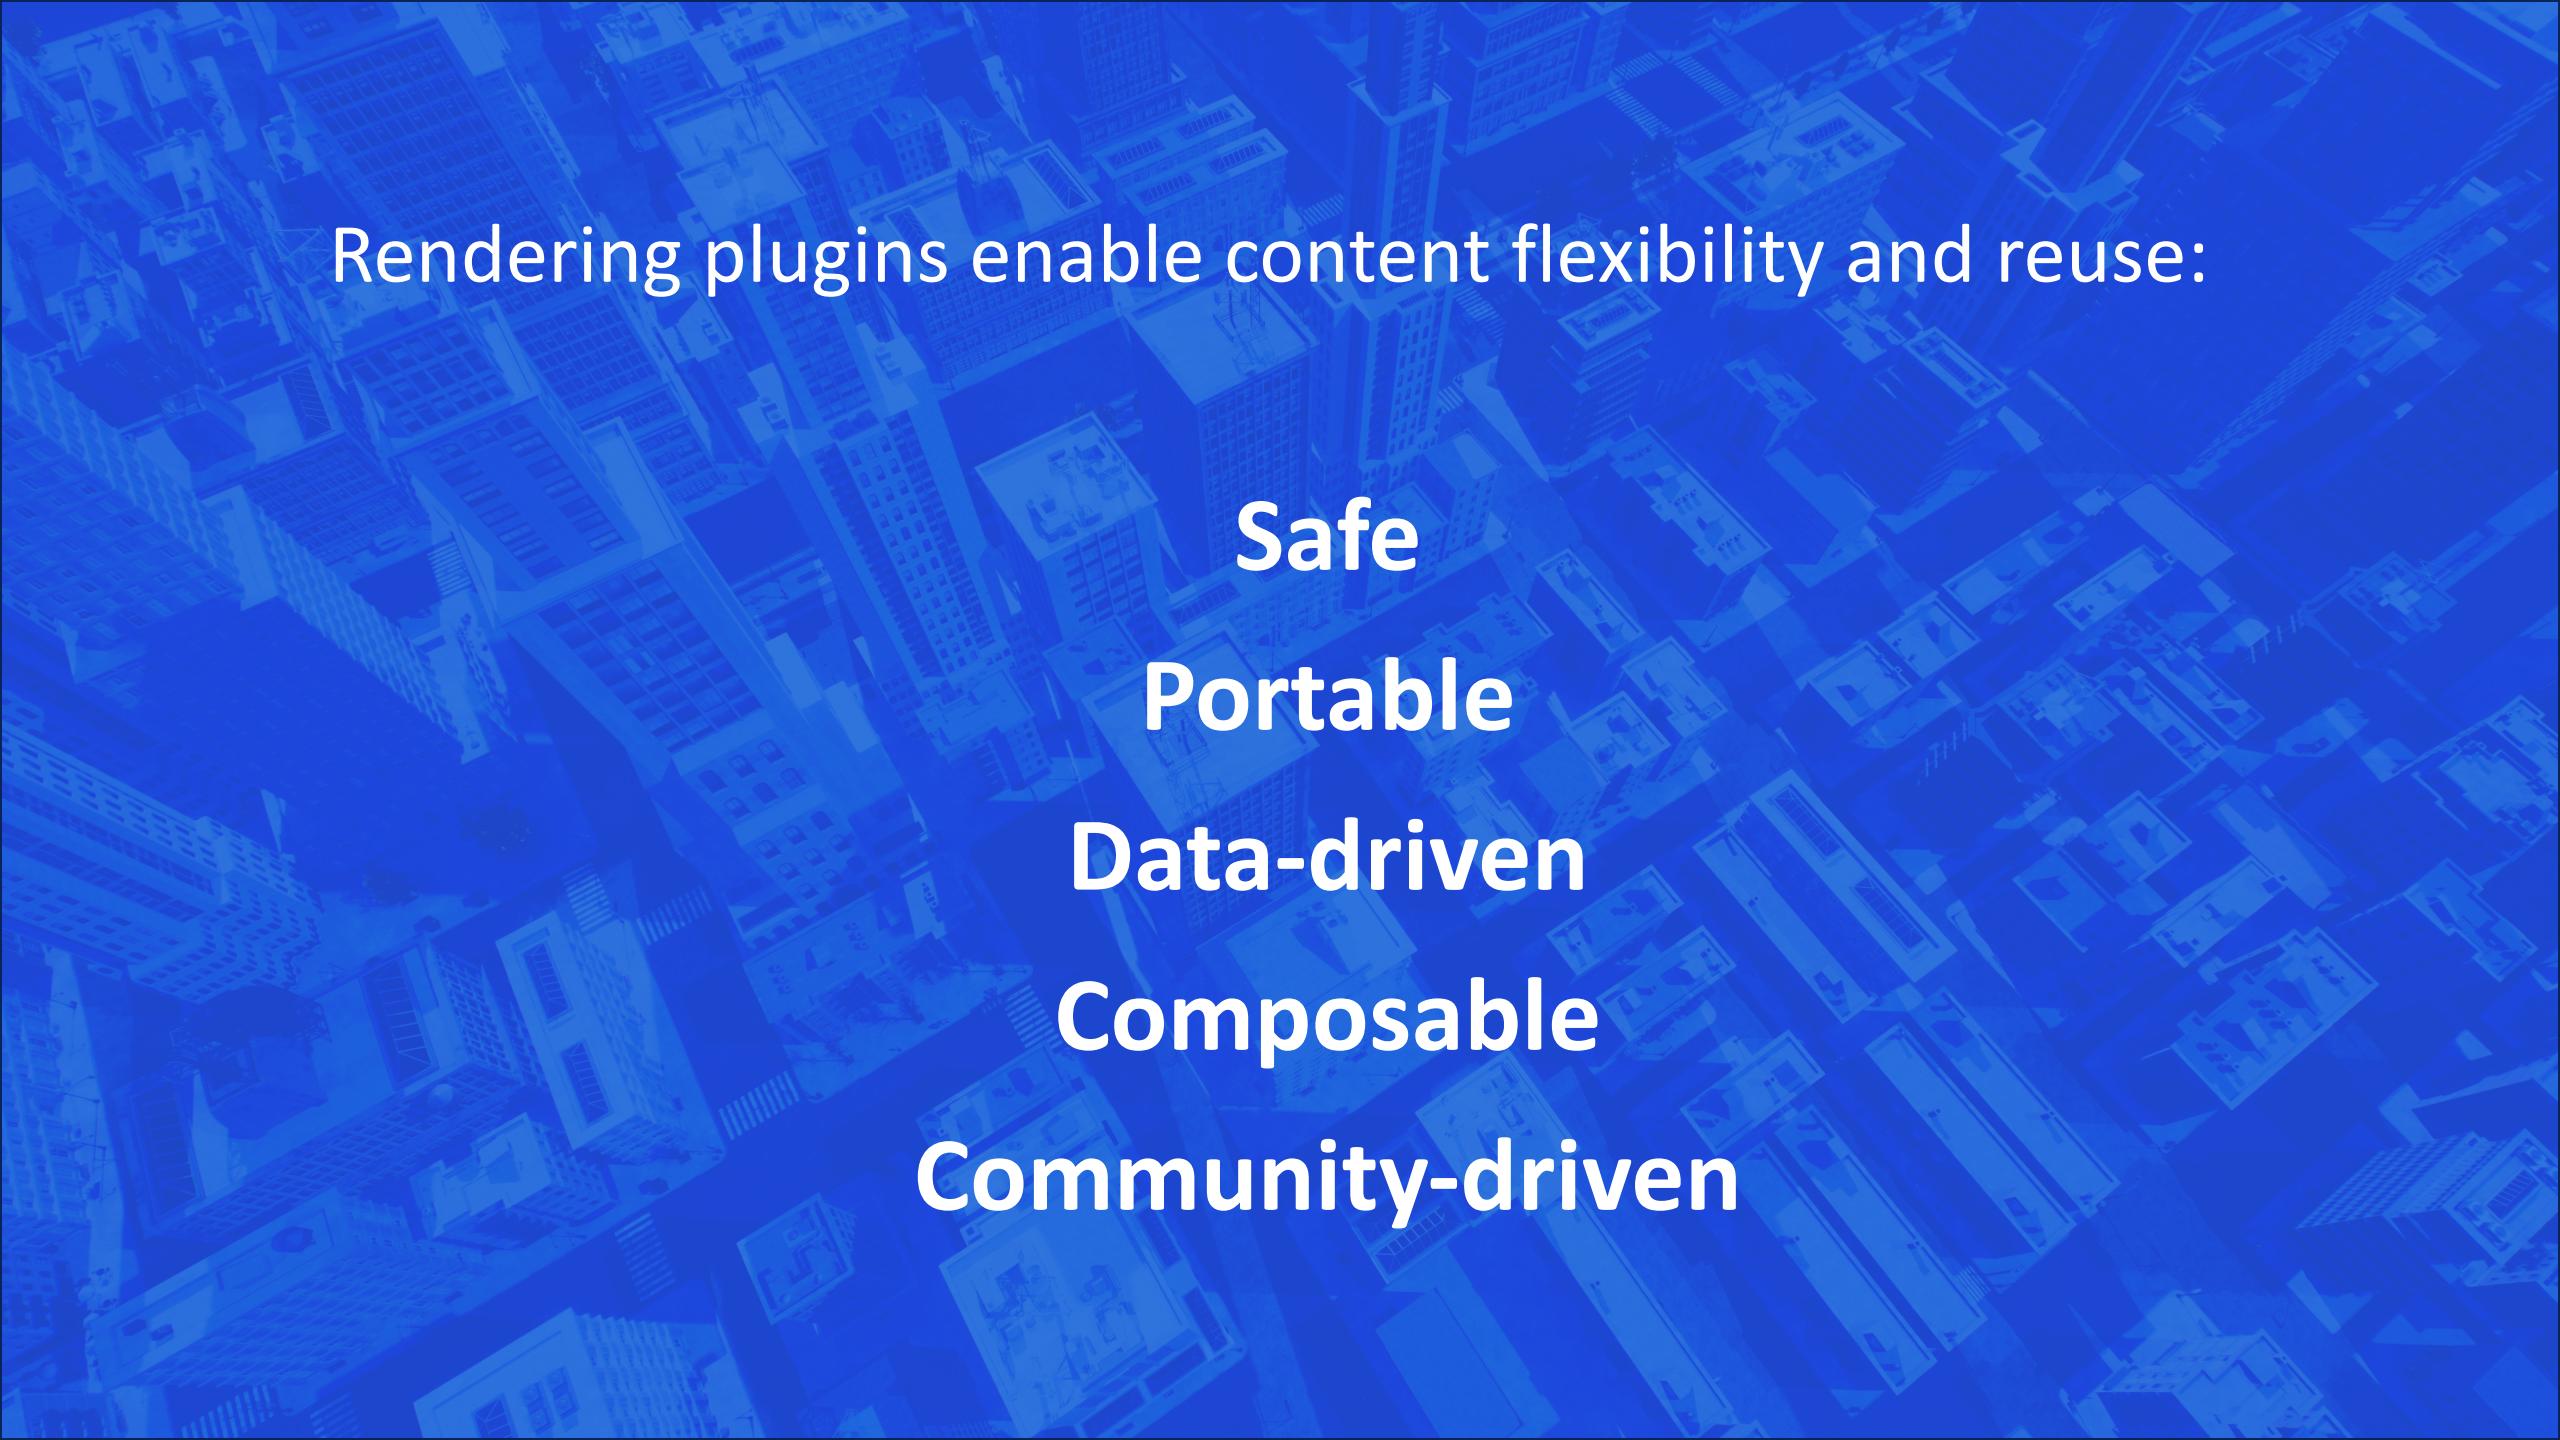


Do

- Use automatic mapping systems
- Use concurrency model of host app

Capability and resource restriction

- Capability restriction depends on host
 - “Does this plugin need full WebGPU API”
 - Can it mine bitcoin with arbitrary shaders?
 - Maybe restrict
- Resource restriction depends on frame / object budget
 - CPU time – “Fuel”
 - GPU time – Queries / Indirect?
 - VRAM – “heap / plugin”
 - Object count / format



Rendering plugins enable content flexibility and reuse:

Safe
Portable
Data-driven
Composable
Community-driven



Interested in plugins?
Come chat with us!

Sign up for updates: renderlet.com

Follow along on Twitter: @_renderlet

sean@renderlet.com